

Week_8_Project

Arun Prasad Patel

Nepal Stock Exchange (NEPSE)-Based Projects NEPSE Stock Price Prediction

Objective: Predict the future prices of stocks listed on NEPSE.

Data: Historical stock prices (available from NEPSE website or APIs).

Tasks: Normalize prices, handle missing data, and use Linear Regression or RandomForest Regressor to predict stock prices

Data was collected from nepse alpha site

- Daily nepse price data was collected from 2019 Dec to 2024 Dec
- It contains dayHigh price, Low price, Open price, Close price, date, change percentage and volume

Four models are used to predict the future closing price

- Moving Average
- Linear Regression Model
- Multiple Regression Model
- RandomForest Regressor Model

```
In [55]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as mse
from sklearn.metrics import r2_score
from sklearn.linear_model import LinearRegression
data =
pd.read_csv("data.csv")
data.head(
)
```

Out[55]:

	Symbol	Date	Open	High	Low	Close	Percent Change	Volume
0	NEPSE	2024-12-05	2752.38	2764.29	2717.34	2734.93	-0.57%	8,369,687,402.72
1	NEPSE	2024-12-04	2795.98	2795.78	2747.65	2750.87	-0.89%	9,132,468,906.77
2	NEPSE	2024-12-03	2772.74	2793.03	2751.31	2775.85	0.64%	8,521,219,299.67
3	NEPSE	2024-12-02	2749.27	2773.47	2746.47	2758.04	0.72%	10,541,746,564.07
4	NEPSE	2024-12-01	2762.30	2764.85	2707.94	2738.06	-0.36%	8,832,225,892.82

Exploratory Data Analysis

- CandleStickpatternof recent 100days

```
In [56]: import mplfinance as mpf
```

```
data =  
data.head(100)  
data['Date'] =  
pd.to_datetime(data['Date'])  
data['Volume'] = data['Volume'].str.replace(',', '').astype(float)  
data.set_index('Date', inplace=True) # Set 'Date' as index  
data = data[['Open', 'High', 'Low', 'Close', 'Volume']]  
# Plot candlestick chart using mplfinance  
fig, axes =  
mpf.plot(  
    data  
    ,  
    type='candle'  
    ,  
    volume=True  
    ,  
    title="NEPSE Candlestick Chart of 100 days",  
    style='yahoo'  
    ,  
    ylabel="Price"  
    ,  
    ylabel_lower="Volume"  
    ,  
    returnfig=True  
    ,  
    warn_too_much_data=100  
)  
  
# Invert the x-axis (dates)  
axes[0].invert_xaxis() # Invert the x-axis of the first axis (candlestick chart)  
  
# Show the plot  
plt.show()  
)
```

NEPSE Candlestick Chart of 100 days



Changing date into date time format

```
In [57]: data = pd.read_csv('data.csv')
data["Date"] =
pd.to_datetime(data['Date'])
data.info(
)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1128 entries, 0 to 1127
Data columns (total 8 columns):
 # Column          Non-Null Count  Dtype
---  -
0 Symbol           1128 non-null   object
1 Date             1128 non-null   datetime64[ns]
2 Open             1128 non-null   float64
3 High             1128 non-null   float64
4 Low              1128 non-null   float64
5 Close            1128 non-null   float64
6 Percent Change    1128 non-null   object
7 Volume           1128 non-null   object
dtypes: datetime64[ns](1), float64(4), object(3)
memory usage: 70.6+ KB
```

```
In [58]: data = pd.DataFrame(data)
data.describe(
)
```

	Date	Open	High	Low	Close
count	1128	1128.000000	1128.000000	1128.000000	1128.000000
mean	2022-07-02 20:06:22.978723584	2165.339071	2184.024947	2142.533327	2161.999140
min	2019-12-0800:00:00	1131.920000	1135.370000	1125.810000	1135.370000
25%	2021-05-0218:00:00	1908.567500	1926.862500	1888.872500	1907.935000
50%	2022-07-0312:00:00	2078.000000	2093.210000	2059.245000	2072.005000
75%	2023-09-0512:00:00	2596.872500	2620.352500	2548.730000	2585.002500
max	2024-12-0500:00:00	3208.530800	3227.110000	3178.250000	3198.190000
std	NaN	473.228330	476.801759	465.756438	470.101723

MovingAverage

- It is a technical tool that investors use to determine the trend and price of stocks in the future.
- It uses previous days' average price to determine the price of the next days.
- A 7-day window size is taken here to determine the next day's price.

```
In [59]: import plotly.graph_objects as go

window_size = 7
data['Moving_Avg'] = data['Close'].rolling(window=window_size).mean()
# Create traces for the plot
trace_original = go.Scatter(
    x=data['Date'],
    y=data['Close'],
    mode='lines',
    name='Original'
)
trace_moving_avg = go.Scatter(
    x=data['Date'],
    y=data['Moving_Avg'],
    mode='lines',
    name=f'{window_size}-Day Moving Average'
)
# Layout for the plot
layout = go.Layout(
    title=f'{window_size}-Day Moving Average of Closing Prices',
    xaxis=dict(title='Date'),
    yaxis=dict(title='Close Price')
)
```

```
fig= go. Figure( data=[ trace_original, trace_moving_avg],
layout=layout)
fig.show(
)
valid_indices=
~np. isnan( data[ 'Moving_Avg'] )
actual= data.loc[ valid_indices,
'Close']
moving_avg= data.loc[ valid_indices,
'Moving_Avg']. values
```

```
# Calculate mean squared error
mse= mse( actual, moving_avg)
print(f'Mean Squared Error of the Moving Average: {mse} ')
rmse_value=
np. sqrt( mse)
r2_value= r2_score( actual, moving_avg)

print(f"Root Mean Squared Error (RMSE): {rmse_value} ")
print(f"R-squared (R²):
{r2_value} ")
```

Mean Squared Error of the Moving Average: 2530.8692685143988
Root Mean Squared Error (RMSE): 50.30774561152983
R-squared (R²): 0.9885020365752958

```
In [60]: total_number_of_days= (data. Date. max() -
data. Date. min() ). days
print("total number of days=", total_number_of_days)
```

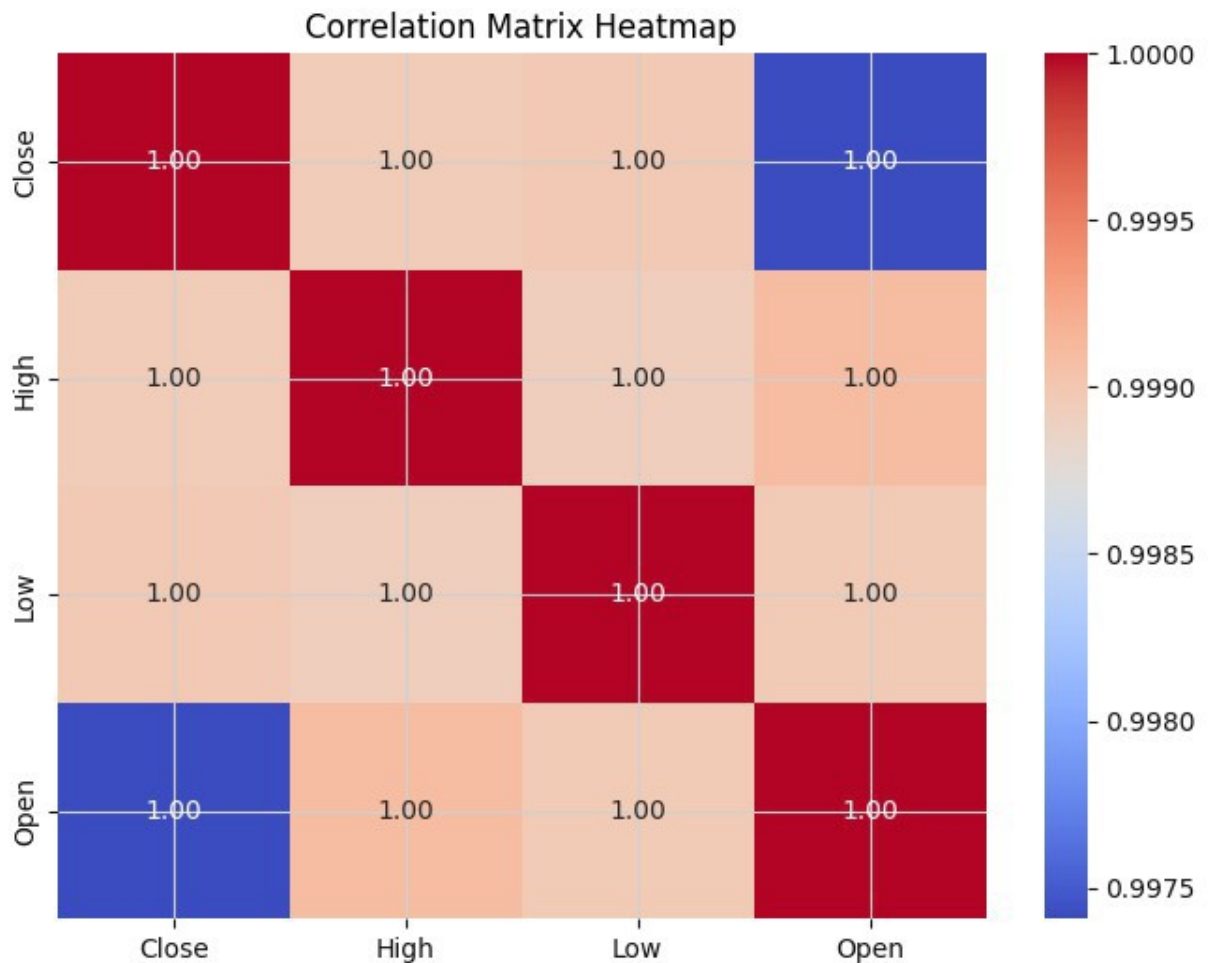
total number of days= 1824

Here, Correlation between closing price, opening price, high price and low price is mapped using heat map

```
In [61]: import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
scaler= MinMaxScaler()
scaler. fit( data[ ['Close']]
)
data[ 'Close'] =
scaler. transform( data[ ['Close']] )
scaler. fit( data[ ['High']]
)
data[ 'High'] =
scaler. transform( data[ ['High']] )
scaler. fit( data[ ['Open']]
)
data[ 'Open'] =
scaler. transform( data[ ['Open']] )

print( data[ ["Close", 'High', "Low", 'Open']]. corr(
))
plt. figure( figsize=(8,6)) # Set figure
size
sns. heatmap( data[ ["Close", 'High', "Low", 'Open']]. corr(), annot=True,
cmap='coolwarm'
plt. title( "Correlation Matrix
Heatmap")
plt. show(
)
```

	Close	High	Low	Open
Close	1.000000	0.998953	0.998977	0.997406
High	0.998953	1.000000	0.998919	0.999096
Low	0.998977	0.998919	1.000000	0.998959
Open	0.997406	0.999096	0.998959	1.000000



This shows that there is a high correlation between the high, low, close and open price.

Linear Regression Model

- Linear Regression model has equation $y = mx + c$
 - y : The predicted value (the scaled 'Close' values).
 - x : The independent variable (the day index).
 - m : slope of the linear model
 - c : the intercept
- Data is scaled between 0 and 1 to improve the model performance
- Split into training (80%) and testing (20%) sets.
- A linear regression model is trained on the training data.
- Predictions are made and compared against the actual test data.
- Performance metrics (MSE, RMSE, R^2) evaluate the model.
- A visualization contrasts the actual vs. predicted values for training data.

```
In [62]: from sklearn.metrics import mean_squared_error as mse
```

```
data = pd.read_csv('data.csv')
```

```

scaler=MinMaxScaler()
scaler.fit(data[['Close']])
data['Close']=
scaler.transform(data[['Close']])

X=np.array(data.index).reshape(-1,
1)
y=
data['Close']
X_train,X_test,y_train,y_test=train_test_split(
    X,y,test_size=0.2,shuffle=False,
    random_state=42
)
model=LinearRegression()
model.fit(X_train,
y_train)
y_pred=
model.predict(X_test)
# Trace for actual data points (scatter plot)
trace0=go.Scatter(
    x=X_train.T[0]
],
    y=y_train
,
    mode='markers'
,
    name="actual"
"
)
# Trace for predicted data points (line plot)
trace1=go.Scatter(
    x=X_train.T[0]
],
    y=model.predict(X_train).
T,
    mode='lines'
,
    name="predicted"
"
)
# Combine both traces into a list
predicted_data=[trace0,trace1]
layout=go.Layout(
    title="Actual vs Predicted Data",
    xaxis=dict(title='Day',
autorange='reversed'),
    yaxis=dict(title='Target
• ')
)
plot=go.Figure(data=predicted_data,
layout=layout)
plot.show()
)
mse_value=mse(y_test,y_pred)
rmse_value=
np.sqrt(mse_value)
r2_value=r2_score(y_test,y_pred)

```



```
print(f"Mean Squared Error (MSE): {mse_value}")  
print(f"Root Mean Squared Error (RMSE): {rmse_value}")  
print(f"R-squared (R²):  
{r2_value}")
```

Mean Squared Error (MSE): 0.3109412041107992

Root Mean Squared Error (RMSE): 0.5576210219412456

R-squared (R²): -9.649145799000943

There is a lot of error in this model making it irrelevant for predicting the price.

Multiple Linear Regression model

It uses more than one independent variable for predicting dependent variables.

- It uses High, low, open price to predict the close price.
- The model splits dataset into train (80%) and test (20%).

```
In [63]: import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as mse
from sklearn.metrics import r2_score

scaler = MinMaxScaler()
scaler.fit(data[['Close']])
data['Close'] = scaler.transform(data[['Close']])
scaler.fit(data[['High']])
data['High'] = scaler.transform(data[['High']])
scaler.fit(data[['Open']])
data['Open'] = scaler.transform(data[['Open']])
X = data[['Open', 'High', 'Low']] # Independent variables
y = data['Close'] # Dependent variable


X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

mse_value = mse(y_test, y_pred)
r2_value = r2_score(y_test, y_pred)
model.fit(X, y)
y_pred_full = model.predict(X)

plt.figure(figsize=(12, 6))
plt.plot(data.index, y, color='blue', label='Actual Price')
plt.plot(data.index, y_pred_full, color='red', linestyle='--', label='Predicted Price')
plt.xlabel('Day')
plt.gca().invert_xaxis()
plt.ylabel('Close Price')
plt.title('Actual vs Predicted Close Price for All Days')
plt.legend()
plt.tight_layout()
plt.show()
```

```
plt.show(  
)
```

```
print("Model Coefficients:",  
      model.coef_)  
print("Intercept:",  
      model.intercept_)  
print("Mean Squared Error (MSE):", mse_value)  
print("R-squared:",  
      r2_value)
```



Model Coefficients: [-8.06393657e-01 9.42776311e-01 4.22032540e-04]

Intercept: -0.47788042456870206

Mean Squared Error (MSE): 3.206797429870102e-05

R-squared: 0.9993654540440245

RandomForest Regressor

- RandomForest Regressor uses multiple decision trees to decide the outcome and chooses the best decision tree
- A dataset is divided into train (80%) and test (20%) data.

```
In [64]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split, TimeSeriesSplit
features = ['Open', 'High', 'Low']
target = 'Close'
scaler = MinMaxScaler()
data[features] = scaler.fit_transform(data[features])
data[target] = scaler.fit_transform(data[[target]])
X = data[features]
y = data[target]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=False,
    random_state=42
)

rf_model = RandomForestRegressor(n_estimators=100, max_depth=10, random_state=42)
rf_model.fit(X_train, y_train)

# Predict on test data
y_pred_rf = rf_model.predict(X_test)

# Evaluate the model
mse_value = mse(y_test, y_pred_rf)
```

```
r2_value = r2_score(y_test, y_pred_rf)
```

```
# Print evaluation metrics
```

```

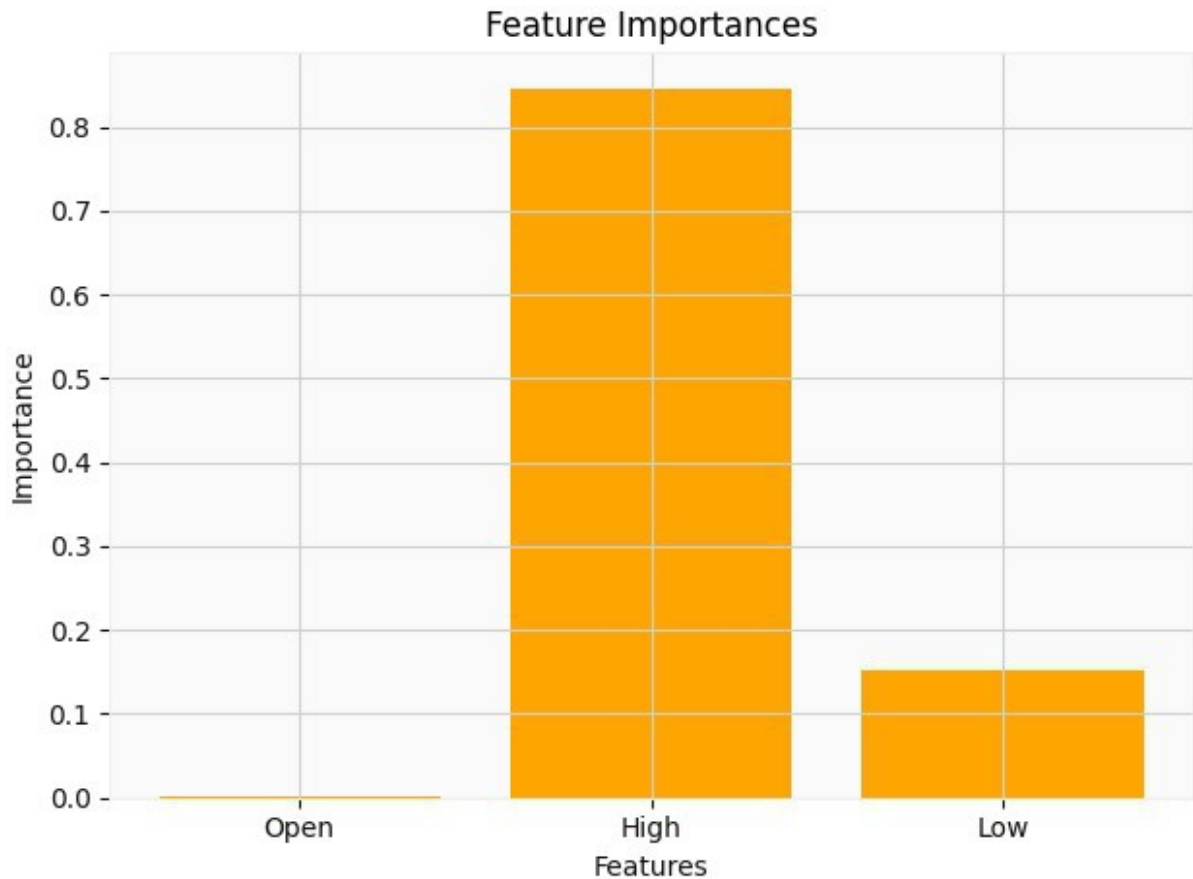
print("Mean Squared Error (MSE):", mse_value)
print("R-squared (R²):", r2_value)

# Predict on full data for visualization
y_pred_full=
rf_model.predict(X)

# Plot actual vs predicted prices
plt.figure(figsize=(12,
6))
plt.plot(data.index,y,label="Actual Price",
color='blue')
plt.plot(data.index,y_pred_full,label="Predicted Price",color='red',
linestyle='
plt.axvline(X_train.index[-1],color='green',linestyle='--',label='Train-Test
Spl
plt.xlabel('Day'
)
plt.gca().invert_xaxis
()
plt.ylabel('Close
Price')
plt.title('Actual vs Predicted Close
Price')
plt.legend(
)
plt.tight_layout(
)
plt.show(
)

```





Conclusion

This Python code serves the purpose of exploring data and developing models through the application of different machine learning approaches. It starts by getting data from a CSV file and preparing it. Exploratory Data Analysis (EDA) involves making candlestick charts and looking at correlation matrices. Four models Moving Average, Linear Regression, Multiple Linear Regression, and Random Forest Regressor are used to predict stock prices. The Moving Average method makes data clearer by averaging it over a 7-day period. Linear Regression and Multiple Regression models help predict prices. We use measures like MSE, RMSE, and R^2 to check how good the predictions are. The Random Forest Regressor is strong because it uses decision trees and helps us understand which features are important. The study looks at how stock prices are related to each other and checks how well each model can predict future prices. However, mistakes in some models show that they may not be very reliable for making accurate predictions. The script shows a complete analysis of stock prices for NEPSE. It focuses on creating models and visualizations to understand market trends and make predictions.